

# Given a bank customer, build a neural network-based classifier that can determine whether they will leave or not in the next 6 months.

Dataset Description: The case study is from an open-source dataset from Kaggle. The dataset contains 10,000 sample points with 14 distinct features such as CustomerId, CreditScore, Geography, Gender, Age, Tenure, Balance, etc. Link to the Kaggle project: <https://www.kaggle.com/barelydedicated/bank-customer-churn-modeling> Perform following steps:

1. Read the dataset.
2. Distinguish the feature and target set and divide the data set into training and test sets.
3. Normalize the train and test data.
4. Initialize and build the model. Identify the points of improvement and implement the same.
5. Print the accuracy score and confusion matrix.

```
1 import pandas as pd
2 import numpy as np
3 import seaborn as sns
4 import matplotlib.pyplot as plt #Importing the libraries
```

```
1 df = pd.read_csv("Churn_Modelling.csv")
```

## Preprocessing.

```
1 df.head()
```

	RowNumber	CustomerId	Surname	CreditScore	Geography	Gender	Age	Tenure	Balance	NumOfProducts	Has
0	1	15634602	Hargrave	619	France	Female	42	2	0.00	1	
1	2	15647311	Hill	608	Spain	Female	41	1	83807.86	1	
2	3	15619304	Onio	502	France	Female	42	8	159660.80	3	
3	4	15701354	Boni	699	France	Female	39	1	0.00	2	
4	5	15737888	Mitchell	850	Spain	Female	43	2	125510.82	1	

```
1 df.shape
```

(10000, 14)

```
1 df.describe()
```

	RowNumber	CustomerId	CreditScore	Age	Tenure	Balance	NumOfProducts	HasCrCard
count	10000.00000	1.000000e+04	10000.000000	10000.000000	10000.000000	10000.000000	10000.000000	10000.000000
mean	5000.50000	1.569094e+07	650.528800	38.921800	5.012800	76485.889288	1.530200	0.625800
std	2886.89568	7.193619e+04	96.653299	10.487806	2.892174	62397.405202	0.581654	0.486600
min	1.00000	1.556570e+07	350.000000	18.000000	0.000000	0.000000	1.000000	0.000000
25%	2500.75000	1.562853e+07	584.000000	32.000000	3.000000	0.000000	1.000000	0.000000

1 df.isnull()

	RowNumber	CustomerId	Surname	CreditScore	Geography	Gender	Age	Tenure	Balance	NumOfProducts	HasCrCard
0	False	False	False	False	False	False	False	False	False	False	False
1	False	False	False	False	False	False	False	False	False	False	False
2	False	False	False	False	False	False	False	False	False	False	False
3	False	False	False	False	False	False	False	False	False	False	False
4	False	False	False	False	False	False	False	False	False	False	False
...	...	...	...	...	...	...	...	...	...	...	...
9995	False	False	False	False	False	False	False	False	False	False	False
9996	False	False	False	False	False	False	False	False	False	False	False
9997	False	False	False	False	False	False	False	False	False	False	False
9998	False	False	False	False	False	False	False	False	False	False	False
9999	False	False	False	False	False	False	False	False	False	False	False

10000 rows × 14 columns



1 df.isnull().sum()

RowNumber	0
CustomerId	0
Surname	0
CreditScore	0
Geography	0
Gender	0
Age	0
Tenure	0
Balance	0
NumOfProducts	0
HasCrCard	0
IsActiveMember	0
EstimatedSalary	0
Exited	0
dtype:	int64

1 df.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 10000 entries, 0 to 9999
Data columns (total 14 columns):
#   Column                Non-Null Count  Dtype
---
```

```

0   RowNumber      10000 non-null int64
1   CustomerId     10000 non-null int64
2   Surname        10000 non-null object
3   CreditScore    10000 non-null int64
4   Geography      10000 non-null object
5   Gender         10000 non-null object
6   Age           10000 non-null int64
7   Tenure        10000 non-null int64
8   Balance       10000 non-null float64
9   NumOfProducts 10000 non-null int64
10  HasCrCard     10000 non-null int64
11  IsActiveMember 10000 non-null int64
12  EstimatedSalary 10000 non-null float64
13  Exited        10000 non-null int64
dtypes: float64(2), int64(9), object(3)
memory usage: 1.1+ MB

```

```
1 df.dtypes
```

```

RowNumber      int64
CustomerId     int64
Surname        object
CreditScore    int64
Geography      object
Gender         object
Age           int64
Tenure        int64
Balance       float64
NumOfProducts int64
HasCrCard     int64
IsActiveMember int64
EstimatedSalary float64
Exited        int64
dtype: object

```

```
1 df.columns
```

```

Index(['RowNumber', 'CustomerId', 'Surname', 'CreditScore', 'Geography',
      'Gender', 'Age', 'Tenure', 'Balance', 'NumOfProducts', 'HasCrCard',
      'IsActiveMember', 'EstimatedSalary', 'Exited'],
      dtype='object')

```

```
1 df = df.drop(['RowNumber', 'Surname', 'CustomerId'], axis= 1) #Dropping the unnecessary columns
```

```
1 df.head()
```

	CreditScore	Geography	Gender	Age	Tenure	Balance	NumOfProducts	HasCrCard	IsActiveMember	EstimatedSalary
0	619	France	Female	42	2	0.00	1	1	1	101356.36
1	608	Spain	Female	41	1	83807.86	1	0	1	113524.72
2	502	France	Female	42	8	159660.80	3	1	0	115966.08
3	699	France	Female	39	1	0.00	2	0	0	9866.67
4	850	Spain	Female	43	2	125510.82	1	1	1	72551.08

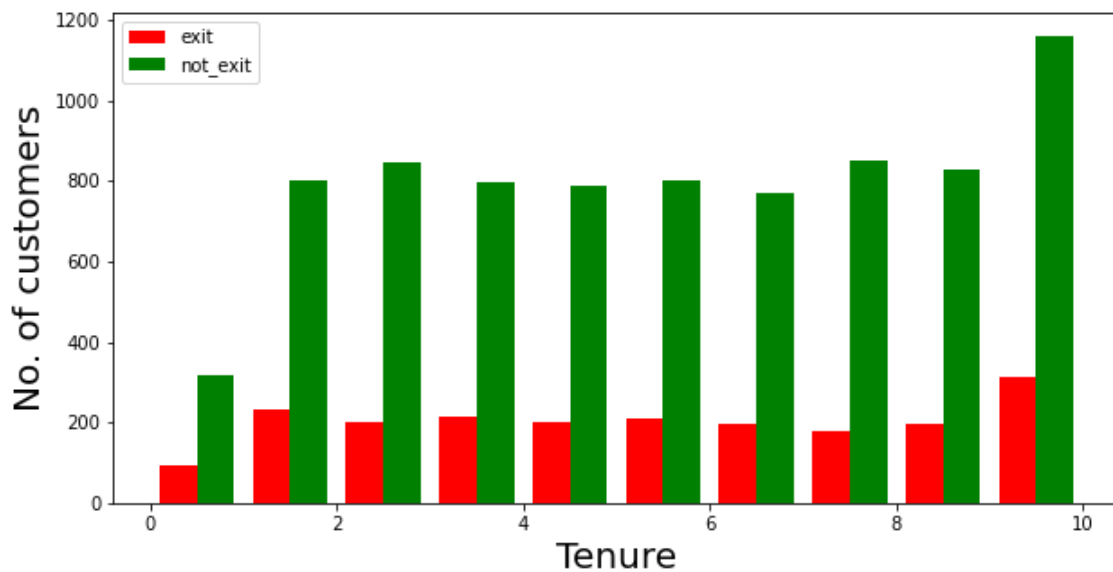
## Visualization

```
1 def visualization(x, y, xlabel):
2     plt.figure(figsize=(10,5))
3     plt.hist([x, y], color=['red', 'green'], label = ['exit', 'not_exit'])
4     plt.xlabel(xlabel,fontsize=20)
5     plt.ylabel("No. of customers", fontsize=20)
6     plt.legend()
```

```
1 df_churn_exited = df[df['Exited']==1]['Tenure']
2 df_churn_not_exited = df[df['Exited']==0]['Tenure']
```

```
1 visualization(df_churn_exited, df_churn_not_exited, "Tenure")
```

```
/usr/local/lib/python3.7/dist-packages/numpy/core/fromnumeric.py:3208: VisibleDeprecationWarning: Creating an
return asarray(a).size
/usr/local/lib/python3.7/dist-packages/matplotlib/cbook/__init__.py:1376: VisibleDeprecationWarning: Creating
X = np.atleast_1d(X.T if isinstance(X, np.ndarray) else np.asarray(X))
```



```
1 df_churn_exited2 = df[df['Exited']==1]['Age']
2 df_churn_not_exited2 = df[df['Exited']==0]['Age']
```

```
1 visualization(df_churn_exited2, df_churn_not_exited2, "Age")
```



## Converting the Categorical Variables

```

1 X = df[['CreditScore', 'Gender', 'Age', 'Tenure', 'Balance', 'NumOfProducts', 'HasCrCard', 'IsActiveMember', 'EstimatedSalary']]
2 states = pd.get_dummies(df['Geography'], drop_first = True)
3 gender = pd.get_dummies(df['Gender'], drop_first = True)

1
2 df = pd.concat([df, gender, states], axis = 1)

```

## Splitting the training and testing Dataset

```
1 df.head()
```

	CreditScore	Geography	Gender	Age	Tenure	Balance	NumOfProducts	HasCrCard	IsActiveMember	EstimatedSalary
0	619	France	Female	42	2	0.00	1	1	1	1013
1	608	Spain	Female	41	1	83807.86	1	0	1	11350
2	502	France	Female	42	8	159660.80	3	1	0	11581
3	699	France	Female	39	1	0.00	2	0	0	9267
4	850	Spain	Female	43	2	125510.82	1	1	1	72549

```
1 X = df[['CreditScore', 'Age', 'Tenure', 'Balance', 'NumOfProducts', 'HasCrCard', 'IsActiveMember', 'EstimatedSalary', 'Exited']]
```

```
1 y = df['Exited']
```

```

1 from sklearn.model_selection import train_test_split
2 X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.30)

```

## Normalizing the values with mean as 0 and Standard Deviation as 1

```

1 from sklearn.preprocessing import StandardScaler
2 sc = StandardScaler()

```

```

1 X_train = sc.fit_transform(X_train)
2 X_test = sc.transform(X_test)

```

```
1 X_train
```

```
array([[ 0.33912606, -1.41493093,  1.38251425, ..., -1.08628092,
        -0.56679212,  1.72941551],
```

```

[-1.58539179,  2.10250049,  0.68783334, ...,  0.92057219,
 1.7643153 , -0.57823004],
[ 0.67022591, -0.84453664, -1.04886894, ..., -1.08628092,
 1.7643153 , -0.57823004],
...,
[ 1.77734102, -0.17907664, -0.00684757, ...,  0.92057219,
 -0.56679212, -0.57823004],
[-0.44723607, -1.03466807, -0.35418802, ...,  0.92057219,
 1.7643153 , -0.57823004],
[-0.77833592,  0.67651478, -0.35418802, ...,  0.92057219,
 1.7643153 , -0.57823004]])

```

1 X\_test

```

array([[ 0.2563511 , -0.36920807, -1.74354985, ...,  0.92057219,
        -0.56679212,  1.72941551],
       [-1.1301295 , -1.31986521,  1.0351738 , ...,  0.92057219,
        1.7643153 , -0.57823004],
       [ 0.93924453,  0.01105478,  1.0351738 , ...,  0.92057219,
        -0.56679212, -0.57823004],
       ...,
       [-1.34741378, -0.36920807,  0.34049289, ...,  0.92057219,
        -0.56679212, -0.57823004],
       [ 1.01167262, -1.2247995 ,  0.34049289, ..., -1.08628092,
        -0.56679212, -0.57823004],
       [ 0.14253553,  0.1061205 , -1.39620939, ..., -1.08628092,
        -0.56679212,  1.72941551]])

```

## ▼ Building the Classifier Model using Keras

```
1 import keras #Keras is the wrapper on the top of tensorflow
```

```
2 #Can use Tensorflow as well but won't be able to understand the errors initially.
```

```
1 from keras.models import Sequential #To create sequential neural network
```

```
2 from keras.layers import Dense #To create hidden layers
```

```
1 classifier = Sequential()
```

```
1 #To add the layers
```

```
2 #Dense helps to construct the neurons
```

```
3 #Input Dimension means we have 11 features
```

```
4 # Units is to create the hidden layers
```

```
5 #Uniform helps to distribute the weight uniformly
```

```
6 classifier.add(Dense(activation = "relu",input_dim = 11,units = 6,kernel_initializer = "uniform"))
```

```
1 classifier.add(Dense(activation = "relu",units = 6,kernel_initializer = "uniform")) #Adding second hidden layer
```

```
1 classifier.add(Dense(activation = "sigmoid",units = 1,kernel_initializer = "uniform")) #Final neuron will be having 1 unit
```

```
1 classifier.compile(optimizer="adam",loss = 'binary_crossentropy',metrics = ['accuracy']) #To compile the Artificial Neural Network
```

```
1 classifier.summary() #3 layers created. 6 neurons in 1st,6neurons in 2nd layer and 1 neuron in last
```

Model: "sequential\_2"

Layer (type)	Output Shape	Param #
dense_6 (Dense)	(None, 6)	72
dense_7 (Dense)	(None, 6)	42
dense_8 (Dense)	(None, 1)	7

=====  
Total params: 121  
Trainable params: 121  
Non-trainable params: 0  
=====

1 classifier.fit(X\_train,y\_train,batch\_size=10,epochs=50) #Fitting the ANN to training dataset

```
Epoch 1/50
700/700 [=====] - 1s 2ms/step - loss: 0.4939 - accuracy: 0.7953
Epoch 2/50
700/700 [=====] - 1s 2ms/step - loss: 0.4256 - accuracy: 0.7973
Epoch 3/50
700/700 [=====] - 1s 1ms/step - loss: 0.4215 - accuracy: 0.7973
Epoch 4/50
700/700 [=====] - 1s 2ms/step - loss: 0.4179 - accuracy: 0.7973
Epoch 5/50
700/700 [=====] - 1s 2ms/step - loss: 0.4144 - accuracy: 0.8163
Epoch 6/50
700/700 [=====] - 1s 2ms/step - loss: 0.4121 - accuracy: 0.8270
Epoch 7/50
700/700 [=====] - 1s 2ms/step - loss: 0.4099 - accuracy: 0.8311
Epoch 8/50
700/700 [=====] - 1s 2ms/step - loss: 0.4081 - accuracy: 0.8313
Epoch 9/50
700/700 [=====] - 1s 2ms/step - loss: 0.4062 - accuracy: 0.8321
Epoch 10/50
700/700 [=====] - 1s 2ms/step - loss: 0.4056 - accuracy: 0.8330
Epoch 11/50
700/700 [=====] - 1s 2ms/step - loss: 0.4042 - accuracy: 0.8357
Epoch 12/50
700/700 [=====] - 1s 2ms/step - loss: 0.4031 - accuracy: 0.8354
Epoch 13/50
700/700 [=====] - 1s 1ms/step - loss: 0.4020 - accuracy: 0.8374
Epoch 14/50
700/700 [=====] - 1s 2ms/step - loss: 0.4018 - accuracy: 0.8369
Epoch 15/50
700/700 [=====] - 1s 1ms/step - loss: 0.4013 - accuracy: 0.8354
Epoch 16/50
700/700 [=====] - 1s 2ms/step - loss: 0.4005 - accuracy: 0.8360
Epoch 17/50
700/700 [=====] - 1s 2ms/step - loss: 0.3999 - accuracy: 0.8379
Epoch 18/50
700/700 [=====] - 1s 2ms/step - loss: 0.3990 - accuracy: 0.8366
Epoch 19/50
700/700 [=====] - 1s 2ms/step - loss: 0.3987 - accuracy: 0.8377
Epoch 20/50
700/700 [=====] - 1s 2ms/step - loss: 0.3990 - accuracy: 0.8369
Epoch 21/50
700/700 [=====] - 1s 2ms/step - loss: 0.3981 - accuracy: 0.8389
Epoch 22/50
700/700 [=====] - 1s 2ms/step - loss: 0.3979 - accuracy: 0.8373
Epoch 23/50
700/700 [=====] - 1s 2ms/step - loss: 0.3977 - accuracy: 0.8381
```

```
Epoch 24/50
700/700 [=====] - 1s 2ms/step - loss: 0.3968 - accuracy: 0.8379
Epoch 25/50
700/700 [=====] - 1s 2ms/step - loss: 0.3972 - accuracy: 0.8376
Epoch 26/50
700/700 [=====] - 1s 2ms/step - loss: 0.3966 - accuracy: 0.8376
Epoch 27/50
700/700 [=====] - 1s 2ms/step - loss: 0.3965 - accuracy: 0.8381
Epoch 28/50
700/700 [=====] - 1s 2ms/step - loss: 0.3964 - accuracy: 0.8384
Epoch 29/50
700/700 [=====] - 1s 2ms/step - loss: 0.3964 - accuracy: 0.8370
```

```
1 y_pred = classifier.predict(X_test)
2 y_pred = (y_pred > 0.5) #Predicting the result
```

```
94/94 [=====] - 0s 1ms/step
```

```
1 from sklearn.metrics import confusion_matrix, accuracy_score, classification_report
```

```
1 cm = confusion_matrix(y_test, y_pred)
```

```
1 cm
```

```
array([[2314,  68],
       [ 443, 175]])
```

```
1 accuracy = accuracy_score(y_test, y_pred)
```

```
1 accuracy
```

```
0.8296666666666667
```

```
1 plt.figure(figsize = (10,7))
2 sns.heatmap(cm,annot = True)
3 plt.xlabel('Predicted')
4 plt.ylabel('Truth')
```



Text(69.0, 0.5, 'Truth')



```
1 print(classification_report(y_test,y_pred))
```

	precision	recall	f1-score	support
0	0.84	0.97	0.90	2382
1	0.72	0.28	0.41	618
accuracy			0.83	3000
macro avg	0.78	0.63	0.65	3000
weighted avg	0.81	0.83	0.80	3000

